

Linear Models- Pseudocode Gradient Descent

```
1 initialize thetas to random values between [-1, 1]
2 previous_error := 0.0
3 current_error = calculate_error( thetas, data)
4 until abs( current_error - previous_error) < epsilon
5     new_thetas = []
6     for j = 0 to m
7         new_thetas[ j] = thetas[ j] -
                                alpha * derivative( j, thetas, data)
8     thetas = new_thetas
9     previous_error = current_error
10    current_error = calculate_error( thetas, data)
```

Comments

1. `epsilon` needs to be surprisingly small...probably around 1E-07.
2. `alpha` can be problematic. Too big and you “jump out of the cup”; too small and it’ll take *forever*, to converge (defined by your `epsilon`). The best thing to do is to make `alpha` *adaptive*. If your error is increasing instead of decreasing that is a sign that your `alpha` is too big, reduce it. So you might start with `alpha = 0.1` and then check to see if error is decreasing or not, if it is *increasing* then divide `alpha` by 10.
3. If you assume this code is wrapped in a function as described in the self-check, then “data” doesn’t need to have any particular implementation...use one that you like. It should probably have `x_o = 1` in it already (you don’t want to do that more than once). You may want to strip off the “y” into a separate List

Math

The pseudocode above works for both linear regression and logistic regression.

It’s worth repeating, exactly, what the formulas are for these. For linear regression, we have:

$$\hat{y} = \theta x$$

and for logistic regression we have:

$$\hat{y} = \frac{1}{1 + e^{-\theta x}}$$

Remember that for linear regression, \hat{y} can be anything but for logistic regression, \hat{y} is constrained to be between 0 and 1, exclusive.

Each has a different implementations for their `calculate_error()` functions (this actually is true *only* if you are using logistic regression for classification. If you're trying to actually use logistic regression as a numerical model, you need to use the other error function but don't worry about that now).

Remember that the Error for Linear Regression is:

$$J(\theta) = \frac{1}{2n} \sum_i (\hat{y}_i - y_i)^2$$

Quite honestly, it doesn't matter if the "2" is in there as it doesn't make one iota of difference for a *minimization* problem if successive values are multiplied by $\frac{1}{2}$ or not. This entire summation is capsulated in the `calculate_error()` function. (Summation is a loop, you loop over the n data points...unless you're pythonic and then you just use a list comprehension and sum at the end).

The Error for Logistic Regression is trickier but again, it's going to be inside the `calculate_error()` function:

$$J(\theta) = -\frac{1}{n} \sum_i y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)$$

Again, the summation just indicates there is a loop in your `calculate_error()` function. If you're using NumPy, I assume you're figured out that that isn't true for you.

Theoretically, the logs should never cause a problem if y is 0, then you should avoid calculating $\log(\hat{y})$. If $(1 - y)$ is 0, the same is true. However....in *practice*, you can have underflow errors and $\log(\hat{y})$ throw an error because \hat{y} will be 0 when it *theoretically* cannot be. You should check for those errors and just make $\log(\hat{y})$ return a nice default value.

Both of these error functions have the same derivative. It's just that " \hat{y} " is a different function (from above):

$$\frac{\partial J}{\partial \theta_j} = \frac{1}{n} \sum_i (\hat{y}_i - y_i) x_{ij}$$

So what does this mean? Let's skip the summation for now...that's just to calculate an average.

For a specific data point, the adjustment to the j -th θ should be the actual error ($\hat{y} - y$) times that data point's contribution to the error which is, the j -th x . Taking in $(1/n)$ and the summation just says, "adjust the j -th θ *by the average of those weighted errors.*"

Raw LaTeX:

$$J(\theta) = \frac{1}{2n} \sum_i (\hat{y}_i - y_i)^2$$

$$J(\theta) = -\frac{1}{n} \sum_i y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)$$

$$\frac{\partial J}{\partial \theta_j} = \frac{1}{n} \sum_i (\hat{y}_i - y_i) x_{ij}$$